

LARGE LANGUAGE MODELS

I. Introduction

Since the 1950s, artificial intelligence (AI) — the idea that machines or software can replicate human intelligence to answer questions and solve problems — has been an area of significant promise and focus.

Given rapid advancements in computing power and the ability to store and process massive amounts of data, AI has now become commonplace in most daily experiences, as seen with smartphones, connected home devices, intelligent driving features (including self-driving cars), chatbots, and even real estate listings.

Large language models (LLMs) complement and enhance AI applications, and they have become accessible to everyone through tools such as OpenAI's ChatGPT and other generative applications.

1. What are Large Language Models (LLMs)?

1.1. Large Language Models definition

A Large Language Model is a [*Deep Learning*](#) algorithm, so it can perform various [*Natural Language Processing*](#) tasks (also known as NLP). They use large transformer language models and are retrained with large data sets. So they are allowed to recognize, translate, and predict or generate texts with different content.

LLMs, also known as [*Neural Networks \(NNs\)*](#), are computing systems inspired by the human brain. They work by using a network of layered nodes, similar to neurons.

Therefore, they are often trained to perform various tasks such as providing, editing faulty code, checking syntax, or can be trained based on disease descriptions to be able to provide temporary diagnosis or can be used to summarize text, documents requested by users... Where large languages are used to serve many different NLP applications, such as translation, ChatBot and finally AI assistant...




Figure 1. LLMs tasks

1.2. What is Natural Language Processing (NLP)?

Natural Language Processing (NLP) is a form of artificial intelligence (AI) that focuses on how computers can interact with users using human language. It is also a sub-specialty of computational linguistics.

It is an interdisciplinary field that combines computer science, linguistics, and artificial intelligence to study other aspects such as the computation of human language.

1.3. What is a Neural Network?

Neural Networks are a subset of machine learning that mimic the structure and function of the human (biological) brain. They are also known as artificial neural networks (ANNs), and consist of interconnected nodes or artificial neurons, structured in layers and weighted connections to transmit and load data. They have multiple layers that form the basis of many deep learning algorithms.

They are designed to learn patterns and relationships from training data, continually improving and adapting. Applying that knowledge to make predictions or decisions. They have the ability to extract meaningful information from complex data that can solve problems, which sets them apart from traditional algorithms.

2. So, what is a transformer model?

One of the most common architectures of large language models, Large Language Models consist of an encoder and a decoder. They process data by encoding the input data, then simultaneously perform mathematical equations to discover relationships between tokens. This helps the computer recognize patterns that a human would see if it were given the same query.

Transformer models work with self-attention mechanisms, which enables the model to learn more quickly than traditional models like long short-term memory models. Self-attention is what enables the transformer model to consider different parts of the sequence, or the entire context of a sentence, to generate predictions.

3. Key components of Large Language Models

Large Language Models are composed of multiple Neural Network layers. Recurrent layers, feedforward layers, embedding layers, and attention layers work in tandem to process the input text and generate output content.

The embedding layer creates embeddings from the input text. This part of the large language model captures the semantic and syntactic meaning of the input, so the model can understand context.

The feedforward layer (FFN) of a large language model is made of up multiple fully connected layers that transform the input embeddings. In so doing, these layers enable the model to glean higher-level abstractions — that is, to understand the user's intent with the text input.

The recurrent layer interprets the words in the input text in sequence. It captures the relationship between words in a sentence.

The attention mechanism enables a language model to focus on single parts of the input text that is relevant to the task at hand. This layer allows the model to generate the most accurate outputs.

There are three main kinds of large language models:

- **Generic or raw language models** predict the next word based on the language in the training data. These language models perform information retrieval tasks.
- **Instruction-tuned language models** are trained to predict responses to the instructions given in the input. This allows them to perform sentiment analysis, or to generate text or code.
- **Dialog-tuned language models** are trained to have a dialog by predicting the next response. Think of chatbots or conversational AI.

4. What is the difference between Large Language Models and Generative AI?

4.1. *What is Generative AI?*

Generative AI is an umbrella term that refers to Artificial Intelligence models that have the capability to generate content. Generative AI can generate text, code, images, video, and music. Examples of Generative AI include Midjourney, DALL-E, and ChatGPT.

4.2. *Comparison between LLMs and Generation AI*

Generative AI is an umbrella term that refers to artificial intelligence models that have the capability to generate content. Generative AI can generate text, code, images, video, and music. Examples of generative AI include Midjourney, DALL-E, and ChatGPT.

Large language models are a type of generative AI that are trained on text and produce textual content. ChatGPT is a popular example of generative text AI.

All large language models are generative AI. [2]

II. How do Large Language Models work?

A large language model is based on a transformer model and works by receiving an input, encoding it, and then decoding it to produce an output prediction. But before a large language model can receive text input and generate an output prediction, it requires training, so that it can fulfill general functions, and fine-tuning, which enables it to perform specific tasks.

1. Training

Large language models are pre-trained using large textual datasets from sites like Wikipedia, GitHub, or others. These datasets consist of trillions of words, and their quality will affect the language model's performance. At this stage, the large language model engages in unsupervised learning, meaning it processes the datasets fed to it without specific instructions. During this process, the LLM's AI algorithm can learn the meaning of words, and of the relationships between words. It also learns to distinguish words based on context. For example, it would learn to understand whether "right" means "correct," or the opposite of "left."

2. Fine-tuning

In order for a large language model to perform a specific task, such as translation, it must be fine-tuned to that particular activity. Fine-tuning optimizes the performance of specific tasks.

3. Prompt-tuning

Fulfills a similar function to fine-tuning, whereby it trains a model to perform a specific task through few-shot prompting, or zero-shot prompting. A prompt is an instruction given to an LLM. Few-shot prompting teaches the model to predict outputs through the use of examples. For instance, in this sentiment analysis exercise, a few-shot prompt would look like this:

Developer code Review: This code is complicated!

Developer sentiment: Negative

Developer code Review: The performance of this code is superb!

Developer sentiment: Positive

The language model would understand, through the semantic meaning of "hideous" and because an opposite example was provided, that the customer sentiment in the second example is "negative"

Alternatively, zero-shot prompting does not use examples to teach the language model how to respond to inputs. Instead, it formulates the question as "The sentiment in 'This plant is so hideous' is...." It clearly indicates which task the language model should perform, but does not provide problem-solving examples.

III. Use cases

Large language models can be used for several purposes:

1. Information Retrieval

Think of Bing or Google. Whenever you use their search feature, you are relying on a large language model to produce information in response to a query. It's able to retrieve information, then summarize and communicate the answer in a conversational style.

2. Sentiment Analysis

As applications of natural language processing, large language models enable companies to analyze the sentiment of textual data.

3. Text Generation

Large language models are behind generative AI, like ChatGPT, and can generate text based on inputs. They can produce an example of text when prompted. For example: "Write me a poem about palm trees in the style of Emily Dickinson."

4. Code Generation

Like text generation, code generation is an application of generative AI. LLMs understand patterns, which enables them to generate code.

5. Chatbots and Conversational AI

Large language models enable customer service chatbots or conversational AI to engage with customers, interpret the meaning of their queries or responses, and offer responses in turn.

IV. Benefits

With a broad range of applications, large language models are exceptionally beneficial for problem-solving since they provide information in a clear, conversational style that is easy for users to understand.

1. Large set of applications

They can be used for language translation, sentence completion, sentiment analysis, question answering, mathematical equations, and more.

2. Always improving

Large language model performance is continually improving because it grows when more data and parameters are added. In other words, the more it learns, the better it gets. What's more, large language models can exhibit what is called "in-context learning." Once an LLM has been pretrained, few-shot prompting enables the model to learn from the prompt without any additional parameters. In this way, it is continually learning.

3. They learn fast

When demonstrating in-context learning, large language models learn quickly because they do not require additional weight, resources, and parameters for training. It is fast in the sense that it doesn't require too many examples.

V. Limitations and challenges of LLMs

Large language models might give us the impression that they understand meaning and can respond to it accurately. However, they remain a technological tool and as such, large language models face a variety of challenges.

1. Hallucinations

A hallucination is when a LLM produces an output that is false, or that does not match the user's intent. For example, claiming that it is human, that it has emotions, or that it is in love with the user. Because large language models predict the next syntactically correct word or phrase, they can't wholly interpret human meaning. The result can sometimes be what is referred to as a "hallucination."

2. Security

Large language models present important security risks when not managed or surveilled properly. They can leak people's private information, participate in phishing scams, and produce spam. Users with malicious intent can reprogram AI to their ideologies or biases, and contribute to the spread of misinformation. The repercussions can be devastating on a global scale.

3. Bias

The data used to train language models will affect the outputs a given model produces. As such, if the data represents a single demographic, or lacks diversity, the outputs produced by the large language model will also lack diversity.

4. Consent

Large language models are trained on trillions of datasets – some of which might not have been obtained consensually. When scraping data from the internet, large language models have been known to ignore copyright licenses, plagiarize written content, and repurpose proprietary content without getting permission from the original owners or artists. When it produces results, there is no way to track data lineage, and often no credit is given to the creators, which can expose users to copyright infringement issues.

They might also scrape personal data, like names of subjects or photographers from the descriptions of photos, which can compromise privacy.[2] LLMs have already run into lawsuits, including a prominent one by Getty Images [3], for violating intellectual property.

5. Scaling

It can be difficult and time- and resource-consuming to scale and maintain large language models.

6. Deployment

Deploying large language models requires deep learning, a transformer model, distributed software and hardware, and overall technical expertise.

VI. Examples of popular large language models

Popular large language models have taken the world by storm. Many have been adopted by people across industries. You've no doubt heard of ChatGPT, a form of generative AI chatbot.

Other popular LLM models include:

1. PaLM

Google's Pathways Language Model (PaLM) is a transformer language model capable of common-sense and arithmetic reasoning, joke explanation, code generation, and translation.

2. BERT

The Bidirectional Encoder Representations from Transformers (BERT) language model was also developed at Google. It is a transformer-based model that can understand natural language and answer questions.

3. XLNet

A permutation language model, XLNet generated output predictions in a random order, which distinguishes it from BERT. It assesses the pattern of tokens encoded and then predicts tokens in random order, instead of a sequential order.

4. GPT

Generative pre-trained transformers are perhaps the best-known large language models. Developed by OpenAI, GPT is a popular foundational model whose numbered iterations are improvements on their predecessors (GPT-3, GPT-4, etc.). It can be fine-tuned to perform specific tasks downstream. Examples of this are EinsteinGPT, developed by Salesforce for CRM, and Bloomberg's BloombergGPT for finance.

VII. Future advancements in large language models

The arrival of ChatGPT has brought large language models to the fore and activated speculation and heated debate on what the future might look like.

As large language models continue to grow and improve their command of natural language, there is much concern regarding what their advancement would do to the job market. It's clear that large language models will develop the ability to replace workers in certain fields.

In the right hands, large language models have the ability to increase productivity and process efficiency, but this has posed ethical questions for its use in human society.

VIII. References

- [1]. [What are large language models \(LLMs\)?](#)
- [2]. Myer, Mike. [“Are Generative AI and Large Language Models the Same Thing?”](#) Quiq, 12 May 2023
- [3]. Sheng, Ellen. [“In generative AI legal Wild West, the courtroom battles are just getting started”](#) CNBC, April 3, 2023
- [4]. [Getty Images Statement](#), Jan 17 2023 (Accessed June 29, 2023)